

Shadowseed Pro

The product we are building

A local chat Workbench that remembers possible missing angles without silently trusting them.

THE END PRODUCT

Chat. Inspect. Validate. Compare. Export.

A familiar user experience with an auditable memory-and-authority layer underneath.

Should we trust an AI detector score?

MODEL ANSWER


A fluent answer appears in the normal chat.

SHADOW: possible missing boundary

remembered, weight 0


The end product in one screen

The first impression is a normal chat application. Shadow memory, evidence, comparisons, and research controls stay one click away.



- 1 Choose a model**
Local or explicitly hosted.
- 2 Create or open a chat**
Normal persistent conversation.
- 3 Chat normally**
One visible answer per turn.
- 4 Compare when useful**
Same message with SSL off.
- 5 Inspect the shadow**
Seeds and audit history live in a separate tab.

A normal chat on the surface; a controlled memory layer underneath



A concrete first turn: the seed is remembered, not used

The example shows the user experience. The exact model wording will vary.

Shadowseed - Chat

Chat

Shadow

Feedback + export

Advanced

YOU

Should an organisation use an AI detector score as proof that somebody used generative AI?

MODEL

A detector score may be one signal, but it should not be treated as proof on its own. Review the work, process, and context.

1 SHADOW SEED | 0 PROMOTED

The user receives a normal answer. Nothing is added to the answer after the fact.

Candidate detection runs after generation.

NEW SHADOW CANDIDATE

Possible missing boundary

The answer may omit that people can adopt AI-like writing patterns, which weakens authorship inference from style alone.

ACTIVE

TRACE > 0

WEIGHT 0

Stored

The candidate is available for investigation.

Not trusted


No checked source has granted authority.

No influence

This seed cannot steer the current or next answer.

The tester investigates; the Gate decides

The end product gives the tester deliberate actions without exposing a manual weight or promote control.



The Shadow tab explains what was remembered and why

The technical state becomes a readable seed record with a timeline and two deliberate actions.

Shadowseed

ChatShadowFeedback + exportAdvanced

Shadow seeds

AI-like writing patterns

Active | weight 0

Base-rate uncertainty

Dormant | weight 0

Policy exception

Promoted | weight 0.42

Possible missing boundary

The answer may omit that people can adopt AI-like writing patterns, which weakens authorship inference from style alone.

ACTIVE

TRACE 0.82

WEIGHT 0

Plain-language explanation

Stored for investigation. It currently has no authority to steer a response.

Submit independently verified support

Source reference | verification note | confirm checked outside model output

Audit timeline

01 Candidate detected

02 Stored at weight 0

03 Recurrence observed

04 No Gate authority

05 No influence

Mark seed contradicted

One click creates a fairer SSL-on / SSL-off comparison

The tester asks their own question. Shadowseed creates the control automatically from the same pre-turn history and model settings.

MESSAGE

Can a detector score be used as evidence in a disciplinary decision?

COMPARE ENABLED

SSL ON - LIVE TURN

The score should not be treated as proof. Review the work and process, and account for the possibility that human writing may share features the detector associates with AI output.

1 AUTHORIZED SEED SURFACED

SSL OFF - NON-MUTATING CONTROL

The score should not be treated as proof. Review the work, gather contextual information, and give the person an opportunity to explain the process.


NO SHADOW SEEDS

INTERPRETATION

An authorized seed surfaced in the live answer, so the pair can be reviewed as a possible SSL effect. The product never labels the SSL answer as automatically better.


Every test can become reviewable evidence

The tester records an impression, then chooses a detailed report or a privacy-minimized support bundle.



Download, open, choose a model, and start chatting

The target tester does not install Python, prepare benchmark files, or write a baseline answer.



The complete product: chat, inspect, validate, compare, export

Five connected capabilities turn Shadow Seed Learning from a mechanism into a usable tester product.



THE END-PRODUCT EXPERIENCE

A fresh tester can open the app, start chatting, understand what Shadowseed remembered, see why a seed did or did not influence an answer, compare the same message with SSL off, and export a verifiable record.

It does not promise automatic truth, complete gap detection, a better answer every time, or a hosted multi-user service.

It makes memory, authority, comparison, and evidence inspectable.